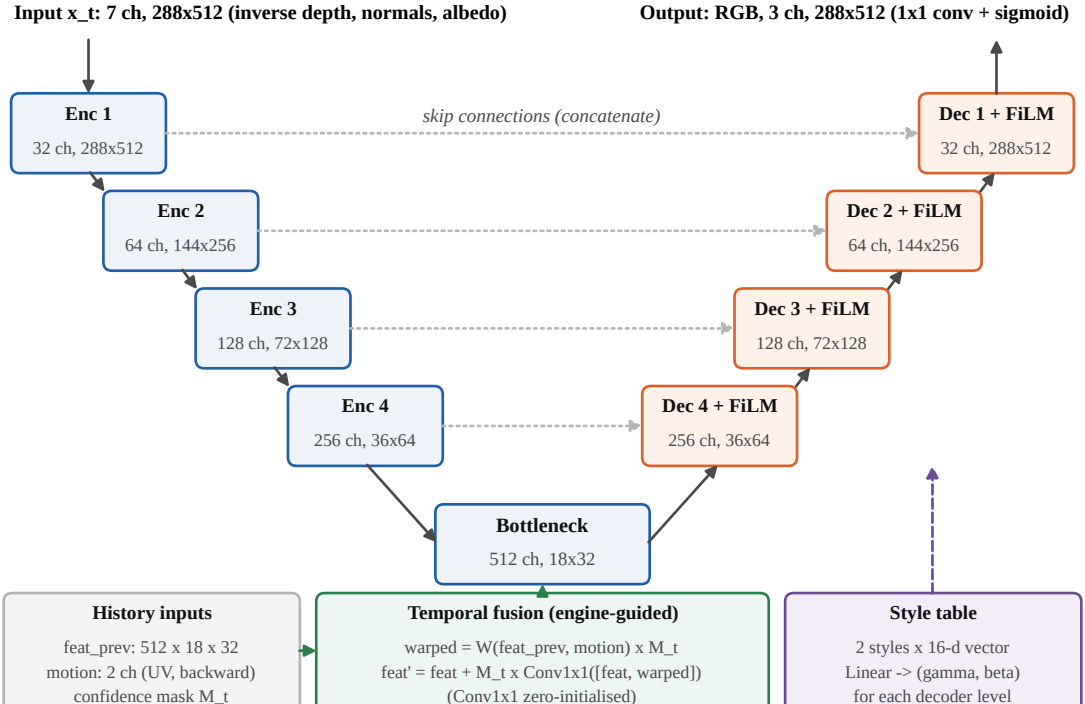
Input x_t: 7 ch, 288x512 (inverse depth, normals, albedo)
Output: RGB, 3 ch, 288x512 (1x1 conv + sigmoid)
Enc 1
32 ch, 288x512
skip connections (concatenate)
Dec 1 + FiLM
32 ch, 288x512
Enc 2
64 ch, 144x256
Dec 2 + FiLM
64 ch, 144x256
Enc 3
128 ch, 72x128
Dec 3 + FiLM
128 ch, 72x128
Enc 4
256 ch, 36x64
Dec 4 + FiLM
256 ch, 36x64
Bottleneck
512 ch, 18x32
History inputs
feat_prev: 512 x 18 x 32
motion: 2 ch (UV, backward)
confidence mask M_t
Temporal fusion (engine-guided)
warped = W(feat_prev, motion) x M_t
feat' = feat + M_t x Conv1x1([feat, warped])
(Conv1x1 zero-initialised)
Style table
2 styles x 16-d vector
Linear -> (gamma, beta)
for each decoder level